

## Question 2 Extra

December 15, 2022

Arunaggiri Pandian Karunanidhi

**0.0.1 I. EXTRA CREDIT (5 points):** Use PyTorch (by switching to a different kernel) to build a simple fully-connected artificial neural network for the beans classification based on the chosen features provided in the data. Generate a confusion matrix for the test data set to demonstrate the accuracy of the model. Based on your model, classify the beans provided in the unlabeled beans-unknown.csv data set. Indicate which classification has been assigned to each of the unlabeled beans. How do the results with the artificial neural network compare to the support vector machine model?

```
In [1]: import pandas as pd
        from sklearn import preprocessing
        import torch
        import torch.nn as nn
        import seaborn as sns
```

```
In [2]: df = pd.read_csv('beans.csv')
        df.head(3)
```

```
Out[2]:
```

	Area	Perimeter	MajorAxisLength	MinorAxisLength	AspectRatio	\
0	28395	610.291	208.178117	173.888747	1.197191	
1	28734	638.018	200.524796	182.734419	1.097356	
2	29380	624.110	212.826130	175.931143	1.209713	

	Eccentricity	ConvexArea	EquivDiameter	Extent	Solidity	roundness	\
0	0.549812	28715	190.141097	0.763923	0.988856	0.958027	
1	0.411785	29172	191.272751	0.783968	0.984986	0.887034	
2	0.562727	29690	193.410904	0.778113	0.989559	0.947849	

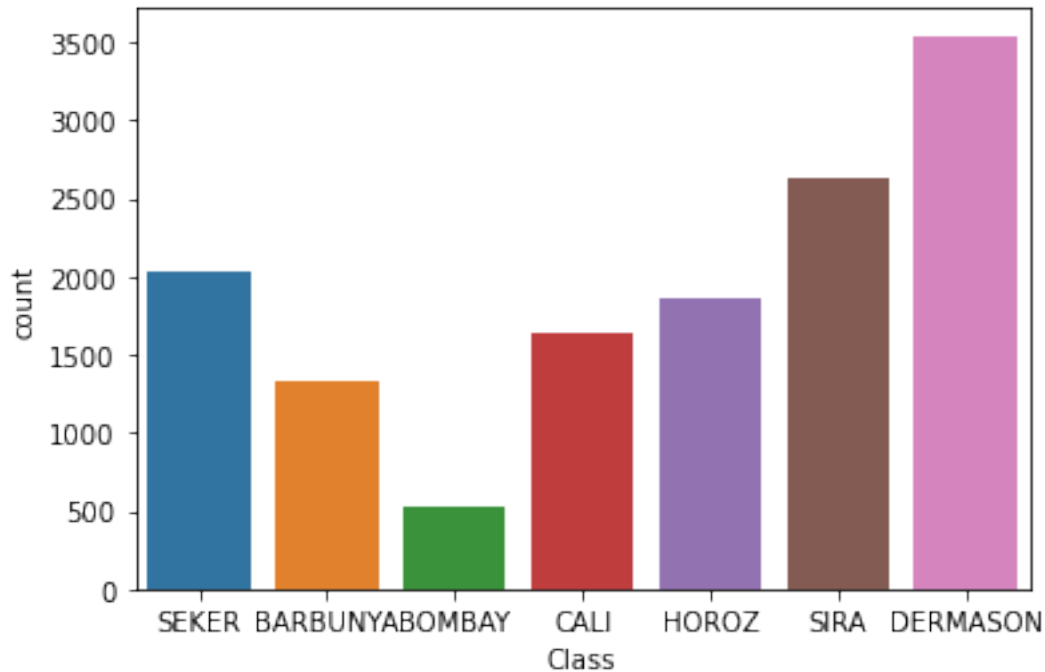
  

	Compactness	ShapeFactor1	ShapeFactor2	ShapeFactor3	ShapeFactor4	Class
0	0.913358	0.007332	0.003147	0.834222	0.998724	SEKER
1	0.953861	0.006979	0.003564	0.909851	0.998430	SEKER
2	0.908774	0.007244	0.003048	0.825871	0.999066	SEKER

```
In [3]: # Visualizing the distribution of the target class
```

```
sns.countplot(x = 'Class', data=df)
```

```
Out[3]: <AxesSubplot:xlabel='Class', ylabel='count'>
```



```
In [4]: df['Class'].unique()
```

```
Out[4]: array(['SEKER', 'BARBUNYA', 'BOMBAY', 'CALI', 'HOROZ', 'SIRA', 'DERMASON'],
              dtype=object)
```

```
In [5]: le = preprocessing.LabelEncoder()
        le.fit(df['Class'])
```

```
Out[5]: LabelEncoder()
```

```
In [6]: df['Class'] = le.transform(df['Class'])
```

```
In [7]: df['Class'].unique()
```

```
Out[7]: array([5, 0, 1, 2, 4, 6, 3])
```

```
In [8]: df.head(3)
```

```
Out[8]:
```

	Area	Perimeter	MajorAxisLength	MinorAxisLength	AspectRatio	\
0	28395	610.291	208.178117	173.888747	1.197191	
1	28734	638.018	200.524796	182.734419	1.097356	
2	29380	624.110	212.826130	175.931143	1.209713	

	Eccentricity	ConvexArea	EquivDiameter	Extent	Solidity	roundness	\
0	0.549812	28715	190.141097	0.763923	0.988856	0.958027	
1	0.411785	29172	191.272751	0.783968	0.984986	0.887034	


2	0.562727	29690	193.410904	0.778113	0.989559	0.947849
	Compactness	ShapeFactor1	ShapeFactor2	ShapeFactor3	ShapeFactor4	Class
0	0.913358	0.007332	0.003147	0.834222	0.998724	5
1	0.953861	0.006979	0.003564	0.909851	0.998430	5
2	0.908774	0.007244	0.003048	0.825871	0.999066	5

```
In [9]: print(df.isnull().sum())
```

```
Area          0
Perimeter     0
MajorAxisLength  0
MinorAxisLength  0
AspectRatio    0
Eccentricity   0
ConvexArea     0
EquivDiameter  0
Extent         0
Solidity       0
roundness      0
Compactness    0
ShapeFactor1   0
ShapeFactor2   0
ShapeFactor3    1
ShapeFactor4    0
Class          0
dtype: int64
```

```
In [10]: sns.countplot(x = 'Class', data=df)
```

```
Out[10]: <AxesSubplot:xlabel='Class', ylabel='count'>
```



In [11]: *# replacing null value with median value*

```
# mode value is
print(df['ShapeFactor3'].median())

# add .mode()[0] as mode returns a series
df['ShapeFactor3'] = df["ShapeFactor3"].fillna(df['ShapeFactor3'].median())

print(len(df[df['ShapeFactor3'].isna()]['ShapeFactor3']))
```

0.6424101875

0

In [12]: *# scaling features*

```
from sklearn.preprocessing import StandardScaler, RobustScaler, MinMaxScaler, PowerTransformer

columns_need_to_be_scaled = df.drop(columns=['Class']).columns

SS = StandardScaler().fit(df[columns_need_to_be_scaled])

df[columns_need_to_be_scaled] = pd.DataFrame(SS.transform(df[columns_need_to_be_scaled]),
                                             columns=columns_need_to_be_scaled)
```

```
df.head(3)
```

```
Out[12]:
```

	Area	Perimeter	MajorAxisLength	MinorAxisLength	AspectRatio	\
0	-0.838853	-1.139688	-1.301921	-0.632217	-1.564982	
1	-0.827322	-1.010590	-1.391089	-0.435922	-1.971943	
2	-0.805349	-1.075346	-1.247768	-0.586894	-1.513942	

	Eccentricity	ConvexArea	EquivDiameter	Extent	Solidity	roundness	\
0	-2.182023	-0.839544	-1.060495	0.288035	0.366749	1.420254	
1	-3.684171	-0.824236	-1.041419	0.697652	-0.465418	0.225234	
2	-2.041468	-0.806883	-1.005376	0.578012	0.517851	1.248936	

	Compactness	ShapeFactor1	ShapeFactor2	ShapeFactor3	ShapeFactor4	Class
0	1.838075	0.681965	2.397588	1.924193	0.838758	5
1	2.496844	0.369690	3.096544	2.690565	0.771240	5
2	1.763526	0.604443	2.230451	1.839561	0.917476	5

```
In [13]: X = df.iloc[:, 0:-1]
```

```
In [14]: y = df.iloc[:, -1]
```

```
In [15]: X = X.to_numpy()
```

```
In [16]: X
```

```
Out[16]: array([[ -0.8388525, -1.13968829, -1.30192147, ...,  2.39758822,
                  1.92419328,  0.83875835],
                [ -0.82732195, -1.01058983, -1.39108908, ...,  3.09654363,
                  2.69056462,  0.77123968],
                [ -0.80534927, -1.07534626, -1.24776821, ...,  2.23045069,
                  1.83956129,  0.91747579],
                ...,
                [ -0.37137187, -0.44579629, -0.44719479, ...,  0.28391186,
                  0.32981927,  0.38873236],
                [ -0.37109977, -0.4250396 , -0.42572569, ...,  0.22306101,
                  0.24219074,  0.03341647],
                [ -0.37069161, -0.38565868, -0.2887106 , ..., -0.1332084 ,
                  -0.28491007,  0.71357655]])
```

```
In [17]: y = y.to_numpy()
```

```
In [18]: y
```

```
Out[18]: array([5, 5, 5, ..., 3, 3, 3])
```

```
In [19]: from sklearn.model_selection import train_test_split
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
```

```

In [20]: # input layer nodes = 16 --> number of features
         # hidden layer nodes = 3
         # output layer nodes = 3 --> number of categories
         iln = 16
         hln = 3
         oln = 7
         eta = 0.01
         num_epoch = 10000

In [21]: # fc = fully connected
         class Net(nn.Module):
             def __init__(self):
                 super(Net, self).__init__()
                 self.fc1 = nn.Linear(iln, hln)
                 self.out = nn.Linear(hln, oln)
             def forward(self, x):
                 x = self.fc1(x)
                 x = nn.functional.relu(x)
                 x = self.out(x)
                 out = nn.functional.softmax(x, dim = 1)
                 return out

In [22]: model = Net()

In [23]: criterion = nn.CrossEntropyLoss()
         optimizer = torch.optim.SGD(model.parameters(), lr = eta)

In [24]: X = torch.Tensor(X_train).float()
         y = torch.Tensor(y_train).long()

In [25]: for epoch in range(num_epoch):
         optimizer.zero_grad()
         out = model(X)
         loss = criterion(out, y)
         loss.backward()
         optimizer.step()
         if epoch % 1000 == 0:
             print('epoch:', epoch, 'loss:', loss.item())

epoch: 0 loss: 1.9576890468597412
epoch: 1000 loss: 1.9133504629135132
epoch: 2000 loss: 1.8884079456329346
epoch: 3000 loss: 1.8597699403762817
epoch: 4000 loss: 1.7924624681472778
epoch: 5000 loss: 1.7611249685287476
epoch: 6000 loss: 1.730120062828064
epoch: 7000 loss: 1.6823604106903076
epoch: 8000 loss: 1.6413053274154663
epoch: 9000 loss: 1.5863434076309204

```

```

In [44]: X = torch.Tensor(X_test).float()
         y = torch.Tensor(y_test).long()

In [45]: out = model(X)

In [46]: (values, prediction) = torch.max(out.data, dim = 1)

In [47]: prediction

Out[47]: tensor([3, 1, 2, ..., 3, 3, 3])

In [48]: prediction = pd.Series(prediction)
         list(prediction)
         len(prediction)

Out[48]: 2707

In [49]: y = list(pd.Series(y))
         len(y)

Out[49]: 2707

In [50]: y = set(y)
         len(y)

Out[50]: 7

In [29]: print('Accuracy is:', (100 * torch.sum(y == prediction).double() / len(y)))

Accuracy is: tensor(73.0329, dtype=torch.float64)

In [67]: test = pd.read_csv('beans-unknown.csv')
         test.head(3)

Out[67]:
```

	Area	Perimeter	MajorAxisLength	MinorAxisLength	AspectRatio	\
0	37500	728.191	275.840463	173.818266	1.586948	
1	37500	715.578	272.171813	175.668301	1.549351	
2	37511	718.350	267.039757	179.141937	1.490660	

	Eccentricity	ConvexArea	EquivDiameter	Extent	Solidity	roundness	\
0	0.776481	37944	218.509686	0.703406	0.988299	0.888690	
1	0.763818	37797	218.509686	0.786229	0.992142	0.920295	
2	0.741599	37868	218.541732	0.717365	0.990573	0.913474	

	Compactness	ShapeFactor1	ShapeFactor2	ShapeFactor3	ShapeFactor4
0	0.792160	0.007356	0.001787	0.627517	0.995836
1	0.802837	0.007258	0.001860	0.644548	0.998631
2	0.818387	0.007119	0.001970	0.669756	0.998379

```
In [68]: # scaling features
```

```
from sklearn.preprocessing import StandardScaler, RobustScaler, MinMaxScaler, PowerTransformer

columns_need_to_be_scaled = test.columns

SS = StandardScaler().fit(test[columns_need_to_be_scaled])

test[columns_need_to_be_scaled] = pd.DataFrame(SS.transform(test[columns_need_to_be_scaled]),
                                              columns=columns_need_to_be_scaled)

test.head(3)
```

```
Out [68]:
```

	Area	Perimeter	MajorAxisLength	MinorAxisLength	AspectRatio	\
0	-1.210167	1.400528	1.669933	-1.592466	1.643560	
1	-1.210167	-1.279132	0.428741	-0.574199	0.490283	
2	0.541390	-0.690215	-1.307557	1.337707	-1.310044	

	Eccentricity	ConvexArea	EquivDiameter	Extent	Solidity	roundness	\
0	1.603313	0.655770	-1.210169	-1.393678	-0.808631	-1.409751	
1	0.527208	-1.639425	-1.210169	0.793939	1.585230	1.272624	
2	-1.360985	-0.530861	0.541415	-1.024995	0.607604	0.693712	

	Compactness	ShapeFactor1	ShapeFactor2	ShapeFactor3	ShapeFactor4
0	-1.658083	1.668114	-1.647227	-1.652531	-1.863595
1	-0.446369	0.441809	-0.455052	-0.451837	0.940218
2	1.318173	-1.299832	1.333936	1.325410	0.688126

```
In [69]: X = test.to_numpy()
```

```
In [70]: X = torch.Tensor(X).float()
```

```
In [71]: out = model(X)
```

```
In [72]: (values, prediction) = torch.max(out.data, dim = 1)
```

```
In [73]: prediction
```

```
Out [73]: tensor([4, 3, 5, 5, 5])
```

```
In [77]: import numpy as np
pre = []
pre = np.array(pre)
for i in prediction:
    if i == 0:
        i = 'BARBUNYA'
    elif i == 1:
        i = 'BOMBAY'
```



```

elif i == 2:
    i = 'CALI'
elif i == 3:
    i = 'DERMASON'
elif i == 4:
    i = 'HOROZ'
elif i == 5:
    i = 'SEKER'
elif i == 6:
    i = 'SIRA'
pre = np.append(pre, [i])

```

In [79]: `len(pre)`

Out[79]: 5

In [82]: `prediction = pd.Series(pre)`

In [83]: `prediction`

```

Out[83]: 0      HOROZ
         1    DERMASON
         2      SEKER
         3      SEKER
         4      SEKER
         dtype: object

```

In [84]: `test['Predicted Class'] = prediction`

In [85]: `test.head(10)`

```

Out[85]:
      Area  Perimeter  MajorAxisLength  MinorAxisLength  AspectRatio  \
0 -1.210167   1.400528         1.669933         -1.592466    1.643560
1 -1.210167  -1.279132          0.428741         -0.574199    0.490283
2  0.541390  -0.690215         -1.307557          1.337707   -1.310044
3  0.859855  -0.333720         -0.444881          0.439992   -0.449276
4  1.019088   0.902540         -0.346237          0.388967   -0.374523

      Eccentricity  ConvexArea  EquivDiameter  Extent  Solidity  roundness  \
0      1.603313    0.655770        -1.210169 -1.393678 -0.808631  -1.409751
1      0.527208   -1.639425        -1.210169  0.793939  1.585230   1.272624
2     -1.360985   -0.530861          0.541415 -1.024995  0.607604   0.693712
3     -0.423925    1.233472          0.859851  0.643799 -1.195077   0.336891
4     -0.345611    0.281044          1.019072  0.980934 -0.189126  -0.893476

      Compactness  ShapeFactor1  ShapeFactor2  ShapeFactor3  ShapeFactor4  \
0     -1.658083    1.668114        -1.647227        -1.652531        -1.863595
1     -0.446369    0.441809        -0.455052        -0.451837         0.940218
2      1.318173   -1.299832          1.333936          1.325410         0.688126

```

3	0.442221	-0.452559	0.434216	0.438891	0.383911
4	0.344058	-0.357531	0.334126	0.340067	-0.148661

	Predicted Class
0	HOROZ
1	DERMASON
2	SEKER
3	SEKER
4	SEKER

### 0.1 How do the results with the artificial neural network compare to the support vector machine model?

The Random Forest Classifier had a better accuracy (94.16%) when compared to the ANN model (73.03%)